KAZO

SIS Sales Intelligence Platform Techno-Commercial Proposal

Automated ingestion, standardisation, validation and reporting of KAZO EBO, LFR and partner sales & stock — the KAZO reporting suite rebuilt as live screens, 15 standard reports, an AI analytics agent (Ask KAZO) and a full Help Center.

Proposal date: 30 September 2026 · Version 1.0 · Validity 30 days

Prepared for: KAZO — Shop-in-Shop / Large-Format-Retail Secondary Sales

Powered by

fundle

1. Executive Summary

KAZO runs ~65 exclusive brand outlets plus large-format-retail and online counters. Sales, stock, transfer and master extracts from Business Central / POS are pulled into a dozen hand-maintained workbooks every week — Broken/Non-Broken sizes, Buying vs Plan, L2L, Product Review, Sell-Through, Offer and Pujo analyses, Footfall — each with its own copy of the data, its own pivots and its own risk of error. Partner (LFR) files add a second set of formats that drift every season.

Fundle proposes the **KAZO SIS Sales Intelligence Platform** — a proven, production-ready web application (screenshots in this document show the platform white-labelled to KAZO; figures shown are illustrative reference data and will be replaced by KAZO's own partner data at onboarding). It ingests every extract and partner file, recognises its structure with a deterministic rules engine (no AI in the data path), validates every row against KAZO masters, routes exceptions to people, and publishes trusted, drill-able reporting — the native **Sales matrix** (hierarchy, slicers, YoY, L2L, targets, productivity), the full **KAZO reporting suite rebuilt as live screens (§3)**, an **Ask KAZO** conversational analytics agent and a Help Center that documents every number.

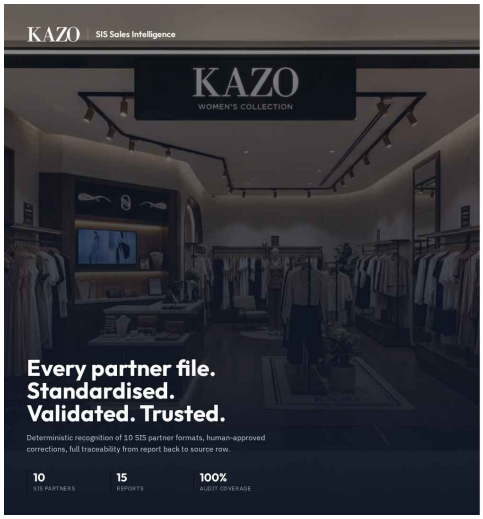
Outcome	What KAZO gets
One source of truth	All partner secondary-sales history (current + last year), every SIS counter, the full style/SKU master, targets, stock and partner margins in a single governed data model
Zero-touch ingestion	Drag-and-drop or email → recognised → validated → published in minutes, with confidence scoring and human approval only where needed
Trust	Every aggregate traces to the batch and source row; financial values are never auto-corrected; complete audit log
Speed of insight	Hero KPIs, 16 reports with filters/sorting/pagination/export, and plain-English questions answered from live data by Ask KAZO
Self-service clarity	Help Center defines every KPI, formula and validation rule exactly as computed

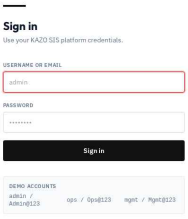
2. Scope Delivered — Modules & Functionality

All modules below are built and running today. Screenshots show the live application in KAZO branding; partner names and figures are illustrative reference data until KAZO's files are onboarded.

2.1 Secure access & role model

Username/e-mail + password login with bcrypt hashing, JWT sessions, brute-force lockout, password change, admin self-lockout protection and automatic administrator recovery. Three roles — **Administrator**, **Data Operations**, **Management** (read-only) — gate every screen and API.

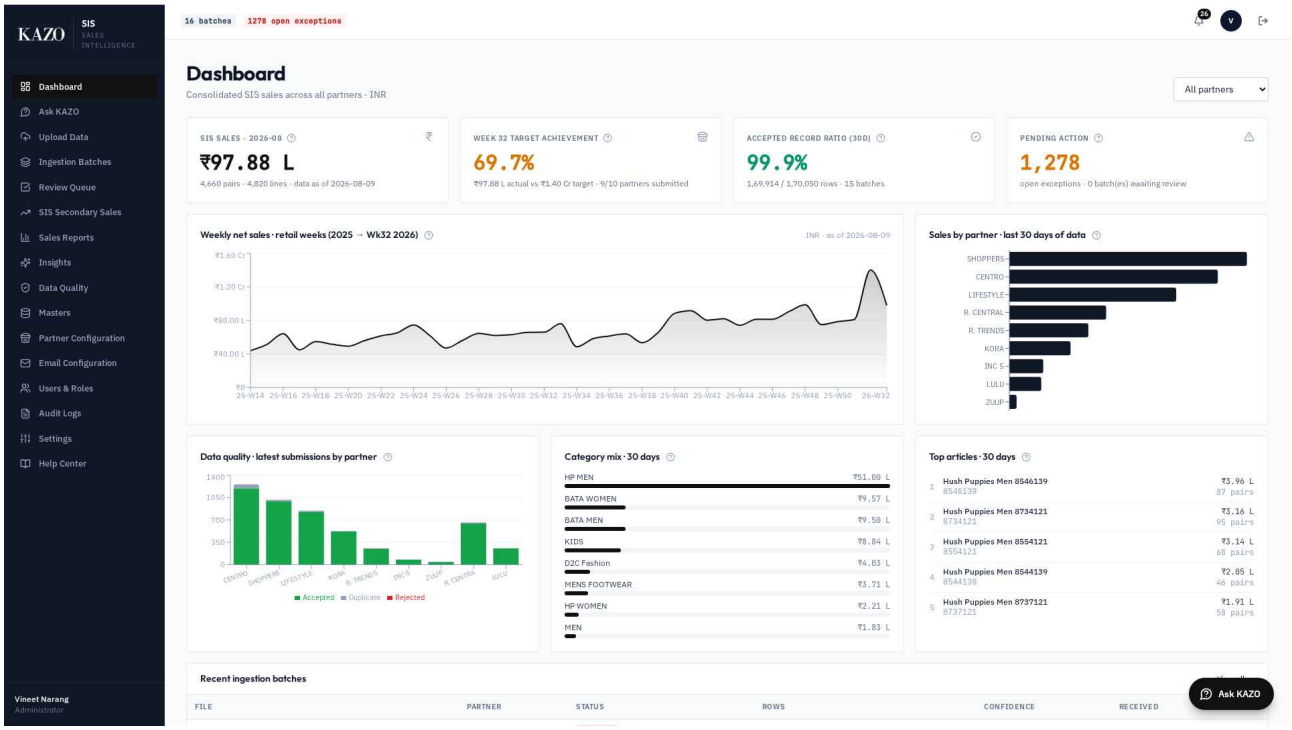




Login — KAZO-branded, role-based entry.

2.2 Executive Dashboard

Month-to-date SIS sales (as of latest data), weekly target achievement, accepted-record ratio, pending actions; weekly trend (last year → current year), sales by partner, category mix, top articles, data-quality by partner and recent batches. Every card carries a “?” that opens the exact formula in the Help Center.

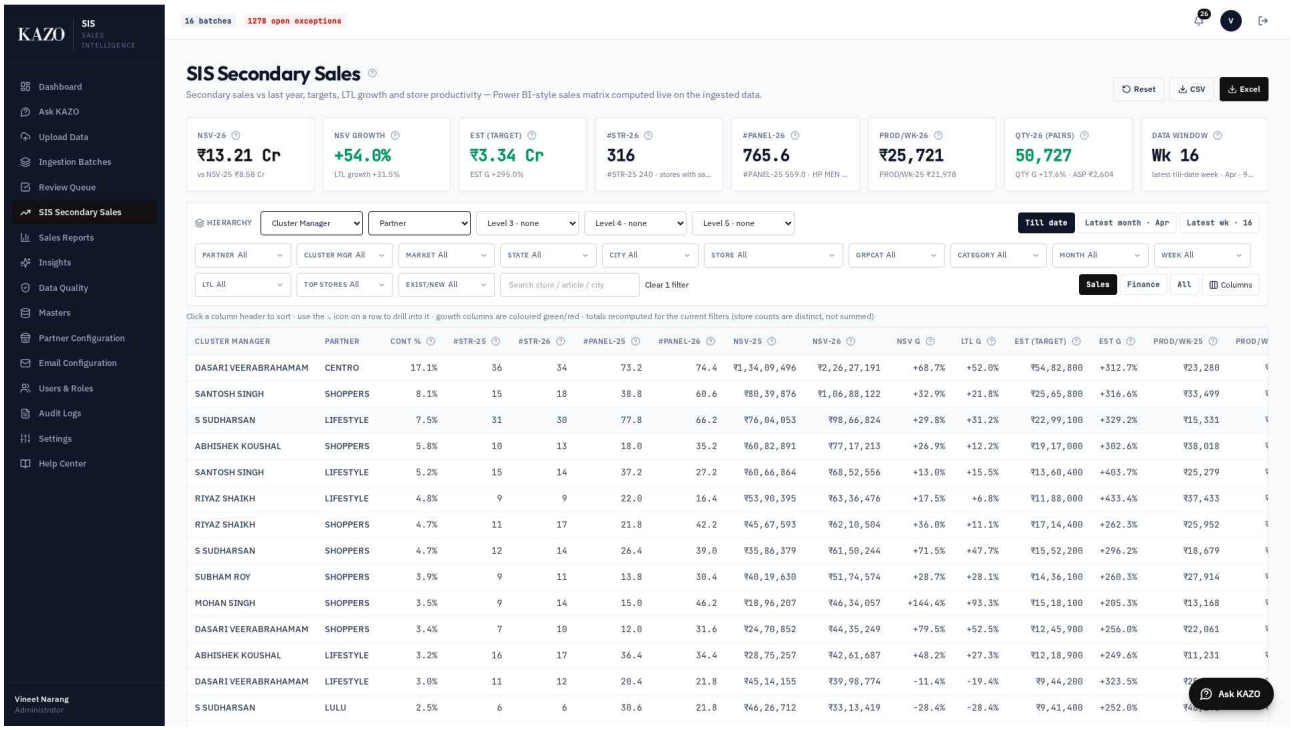


Dashboard — KPIs and charts computed live from the reporting layer.

2.3 Secondary Sales Report (Power BI-style matrix, rebuilt natively)

A Power BI-style sales matrix implemented on the platform's own data model so it never depends on desktop refreshes. The fact table is built from KAZO's partner files and enriched with MRP, cost, partner margin, category/sub-category, cluster/area manager, city/state, LTL (like-to-like) status, top-store and existing/new flags and weekly targets.

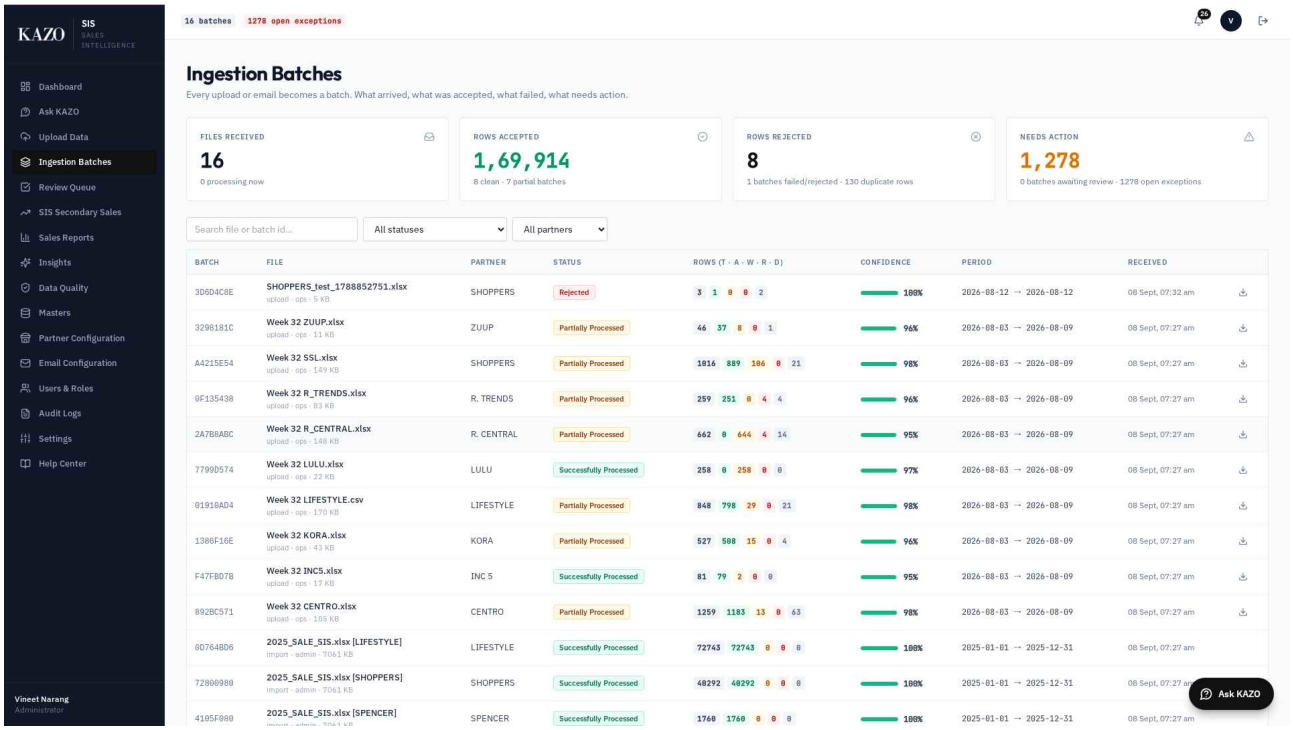
- **Hero bar**: NSV-26, NSV growth, LTL growth, EST target & EST G, #STR, #PANEL, PROD/WK, QTY & ASP, data window.
- **Dynamic hierarchy (5 levels)**: Partner, Area/Cluster Manager, State, Store Code/Name, City, Category, Sub-category, Style, Existing vs New, Top Stores, Month, Week, LTL — any combination, any order; one-click drill-down per row.
- **Slicers**: 13 cascading multi-select slicers + Till Date, Latest Month, Latest Week toggles + free-text search; shareable URL state.
- **Measures (all DAX logic reproduced)**: CONT %, #STR-25/26, #PANEL-LY/CY (configurable panel weights), NSV-25/26, NSV G, LTL G, EST, EST G, STORE PROD/WK-25/26, QTY-25/26, QTY G, ASP; Finance set: MRP VAL, DISC, DISC %, GST VAL, MARGIN VAL, PRIM BILLING, GST REVERSAL, MARGIN REVERSAL, REALIZED VALUE, COST VAL, MARGIN %, PRIM BILLING G, TO G.
- **Enterprise grid**: server-side computation (seconds on 200k+ rows), sortable columns, 50-row pagination, column sets (Sales / Finance / All) and column chooser, green/red growth colouring, distinct-count-aware totals row, CSV & Excel export honouring the current view.



Secondary Sales matrix — hierarchy, slicers, hero KPIs and measure grid.

2.4 Ingestion engine & batches

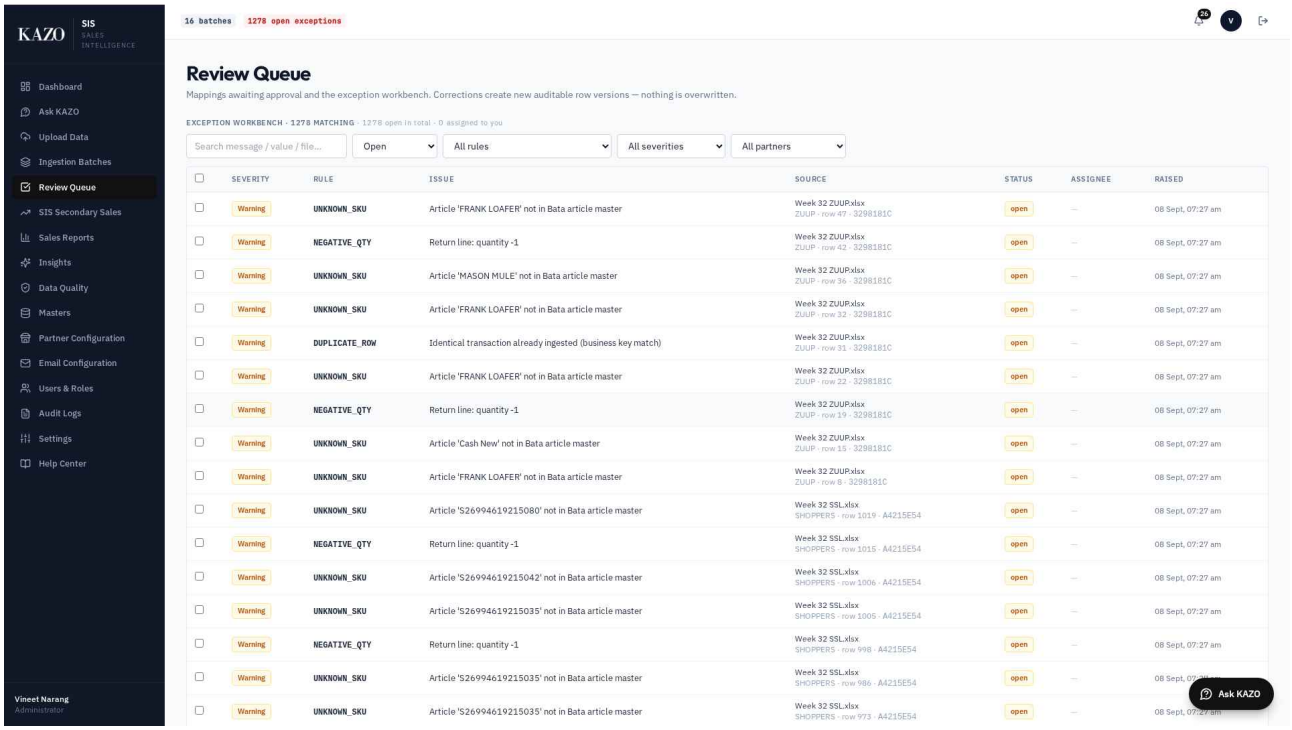
Upload (xlsx/xls/csv, up to 50 MB) or mailbox pickup. Parser detects header rows, skips blank/total rows, coerces numbers/dates. The **recognition engine** identifies partner, header row and column mapping through learned mappings → exact KAZO field → alias dictionary → fuzzy match, and scores confidence (0.3 partner + 0.1 header + 0.6 columns). High confidence auto-accepts; medium/low routes to review. File-hash and business-key duplicate detection; restatements supersede prior batches only after approval; originals preserved in object storage.



Ingestion Batches — hero counters, status, T·A·W·R·D row counts, confidence, download of the original file.

2.5 Validation & Exception Workbench

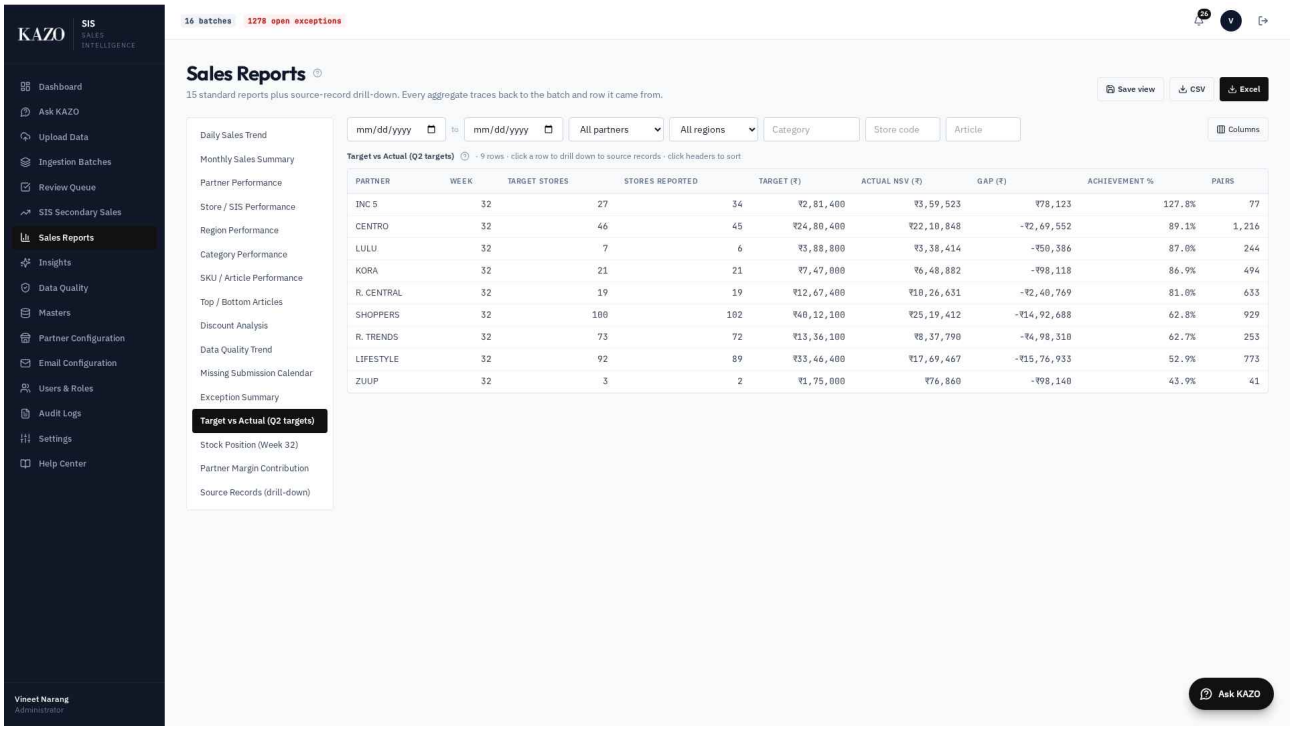
Twelve configurable rules (required field, invalid/future date, unknown store, unknown SKU, negative/zero qty, invalid price, discount over limit, arithmetic mismatch, duplicate row, row-count anomaly, missing submission) with severity and blocking flags. Exceptions are worked by people: accept, correct (versioned), map to master (with remembered alias), add master, ignore, reject row/file, assign, comment, re-validate — individually or in bulk.



Review Queue — mapping approvals and the exception workbench.

2.6 Standard reports & drill-down

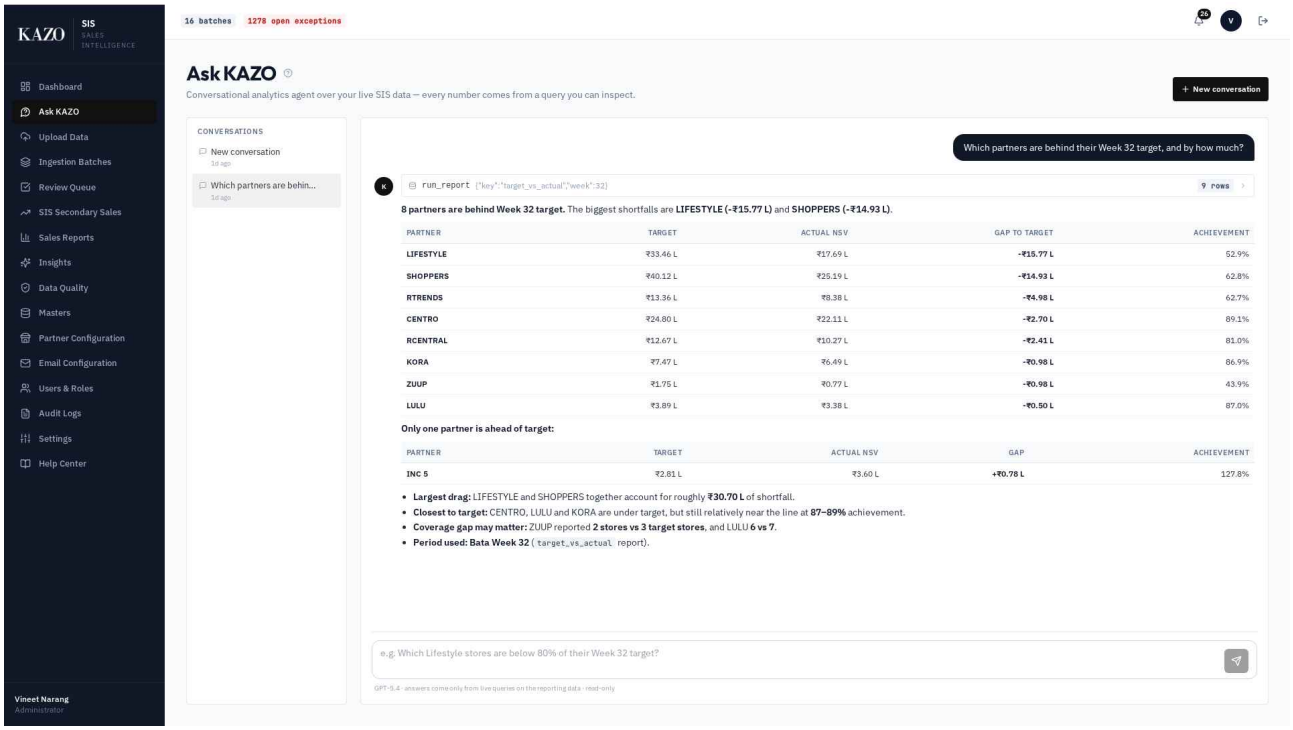
Fifteen reports (daily, monthly, partner, store, region, category, SKU, top/bottom, discount analysis, data-quality trend, missing-submission calendar, exception summary, target vs actual, stock position, partner margin) plus row-level Source Records. Filters, sorting, column selection, saved views, CSV/Excel export; every aggregate row drills to the exact source rows and batch.



Sales Reports — Target vs Actual with drill-down to source records.

2.7 Ask KAZO — AI analytics agent

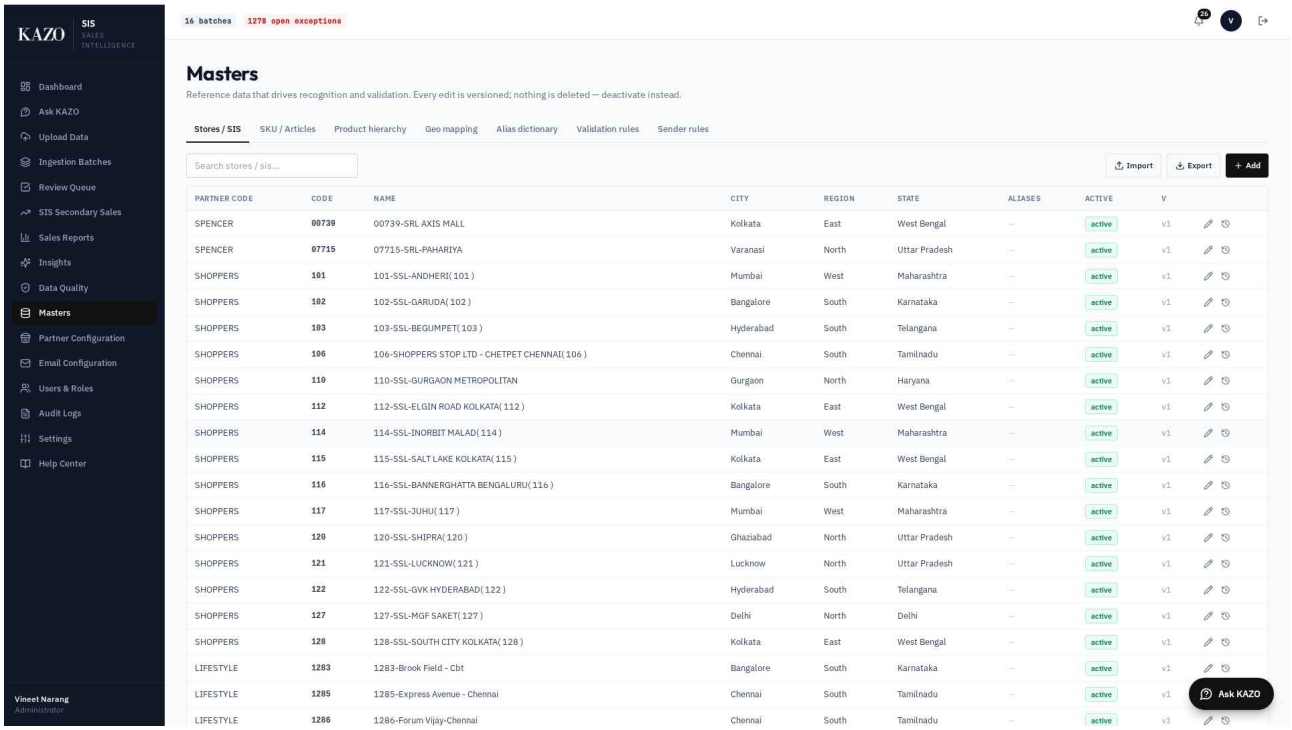
A GPT-5.4 agent grounded in the platform: it never estimates — every figure comes from one of 16 live query tools (reports, ad-hoc aggregation, Secondary Sales measures, store-level targets, stock cover, margins, batches, exceptions, data quality, insights, Help Center). Answers stream with tables and insights; the queries executed are shown and expandable; conversations are saved per user; available as a full page and as a drawer on every screen. Read-only by design — the ingestion engine stays 100% rules-based.



Ask KAZO — question, executed query, tabular answer and analyst commentary.

2.8 Masters & configuration

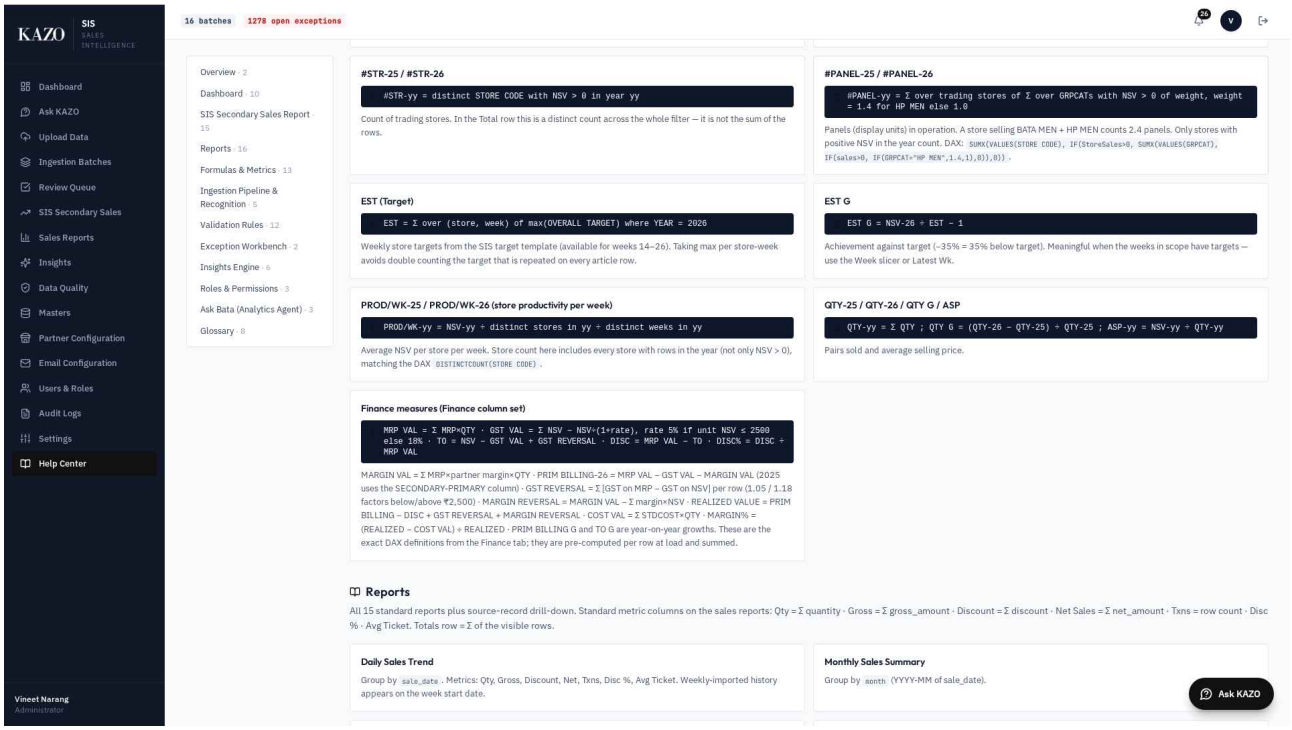
Stores/SIS, articles, product hierarchy, geo mapping, alias dictionary, validation rules, sender rules — versioned, import/export, deactivate-not-delete. Partner configuration (keywords, frequency, partial-ingestion policy), confidence thresholds, notification channels, e-mail connector (Graph/Gmail/IMAP, encrypted credentials), users & roles.



Masters — versioned reference data that drives recognition and validation.

2.9 Help Center

Twelve sections / 100+ topics: every dashboard KPI, every report, all formulas (incl. every Secondary Sales measure), recognition confidence, validation rules, exception actions, insight rules, roles and glossary. Searchable, deep-linked from every “?” in the application, and used verbatim by Ask KAZO so explanations always match.



Help Center — formulas exactly as computed by the platform.

2.10 Also included

- Insights & anomalies: statistical spike/drop, underperforming stores, missing submissions, discount outliers, row-count anomalies, rejection-rate alerts — each linked to records.
- Data Quality page: acceptance %, confidence and exceptions by partner and rule, 30-day trend.
- Notifications (in-app) and complete Audit Log with filters/search.
- White-label KAZO branding (logo, palette, typography), OpenGraph/WhatsApp link previews, responsive desktop/tablet layout, data-testid coverage for automated QA.

3. KAZO Reporting Suite — analysed and in scope

Fundle reviewed the 12 working files shared by KAZO (Kazo Files.zip — Report folder and Offer Analysis sub-folder). Today these are hand-built Excel/XLSB workbooks refreshed from Business Central / POS extracts, each carrying its own copy of the sales dump, stock, item and store masters. The platform replaces them with one governed data model and rebuilds every report as a live, filterable screen with hero KPIs, drill-down, pagination and Excel export — the same engine already shown in §2.3. Every formula is documented in the Help Center.

3.1 Source data identified in the files

Source	Content seen in the files	Platform treatment
POS / BC sales dump	Bill-level lines: store, ASM, zone, date, receipt & invoice, customer, season, occasion, category, sub-category, HSN, SKU/option/variant, colour, size, qty, MRP, discount, gross, tax	Daily automated ingestion → standardised fact `sales_lines`; drives L2L, offers, ATV/UPT, sell-through, product review
Stock (SOH) extract	SKU × location qty with unit price/cost, discount %, MRP value, season group, zone, option & store grade, store MDQ	Daily snapshot → `stock_snapshots`; drives broken-size, ageing, STR %, days cover, size ratio
Transfer orders / in-transit	TO number, posting date, from/to location, qty, value at cost & MRP	→ `transfers`; in-transit added to availability; IST recommendations
Item master (102k SKUs)	SKU, option, item no., variant, colour, size/new size, barcode, price, cost (NAV/BC), brand, main category, category, sub-category, actual/current season, discount %	Versioned master with duplicate/cost-mismatch checks (the file's 'Duplicacy Check' and 'Diff-Nav/BC' columns become validation rules)
Store master (≈65 EBOs + LFR)	Location code, name, status, old/new grade, type, state, zone, cluster/ASM, MDQ, city, L2L flag	Versioned master; grade and MDQ history kept for like-to-like comparisons
Targets & plans	Day-wise store targets, option-wise targets, buy plan (EBO buy qty, re-order qty/date), offer list per option	→ `targets`, `buy_plan`, `offers`; achievement and buying-vs-plan measures
Footfall	Store × month footfall	→ `footfall`; conversion KPI

3.2 Reports to be delivered (one screen each, same filters/sorting/hero-bar pattern)

3.2.1 Broken / Non-Broken Size Report (AW26)

Source: Broken-Non-Broken report 20-Sep AW26.xlsb — 41k SKU-store rows

Store × Option size availability: sizes in stock (XS S M L XL / 26–36), size length, Full-Size vs Broken status, with and without in-transit; cuts by Zone, Store Grade (A+/A/B/C/FO), Option Grade, Category/Sub-category, Season; store count by grade.

Key computations: Broken = option present in store with fewer than the full size set (size table drives Full/Broken); Cut/Full ratio per store & grade; option count with/without in-transit.

3.2.2 Buying vs Plan (AW26)

Source: Buying Vs Plan AW26 20-Sep.xlsb — base data, projections, size ratio, replenishment, in-transit, stock

Option-wise buy plan vs actual: EBO buy qty, in-store date, ageing days & bucket, Store SOH, WH qty, in-transit, STD sale qty / revenue, STR %, ROS, monthly projected sale & closing (Sep–Nov), secured sale, projected STR %, discount budget, target achievement buckets, Top-10 options, size-wise stock ratio vs sale ratio (size gap), replenishment list.

Key computations: STR % = Sale ÷ (Sale + Closing stock) · ROS = Sale qty ÷ selling days · Proj. sale = ROS × days × ageing factor (age-bucket table) · Days cover = SOH ÷ ROS · Size gap = Sale ratio – Stock ratio · Tgt Ach % = STD sale ÷ STD target.

3.2.3 L2L Sale Data (Sep'26)

Source: L2L Sale Data Sep'26.xlsx — 21 sheets: Sep'26 vs Sep'25 sale, opening/closing stock, day-wise & option-wise targets, offer list, wardrobe, in-transit

Like-to-like store performance: Sep'26 vs Sep'25 gross sale, qty, bills by store / zone / ASM; day-wise target vs achievement %; option-wise targets; opening vs closing stock and inward; sale & stock summary by category; live offer list per option.

Key computations: L2L growth = (CY – LY) ÷ LY for stores operational in both periods (Store Master L2L flag) · Day Ach % = actual ÷ day target · Stock movement = Opening + Inward – Sale – Transfers = Closing.

3.2.4 Product Review — Weekly Dashboard

Source: Product Review 20-Sep.xlsb — weekly dashboard, SOH size matrix, IST, aging & grade

Per sub-category weekly review: STD sale qty & gross, option count, Top-5 styles by qty and revenue, ageing buckets (<30, 30–60, 60–90, 90–120, 120+ days) with inventory mix, sales mix and sell-out %; A+/A option list; store × option SOH size matrix; inter-store transfer (IST) suggestions donor → receiver; WH in-transit.

Key computations: Sell-out % = Sale qty ÷ (Sale qty + Inventory) per ageing bucket · Mix % = bucket ÷ total · Option grade from ROS/STR ranking · IST = broken sizes in receiver matched with surplus in donor.

3.2.5 Sell-Through — Category & Store

Source: Sell Through Category Level and Store- 13.Sep.xlsb

Sell-through by category and by store with drop-down selection; stock (SOH) vs sale since season launch (06-Jun'26).

Key computations: Sell Through = Sale qty ÷ (Sale qty + Stock qty).

3.2.6 Offer Analysis — ATV / UPT week-on-week

Source: Offer Analysis ATV 15-Sep & 21-Sep.xlsb

Zone-wise comparison of offer weeks vs previous weeks: bill count, qty, gross sale, discount %, ASP, UPT, discount value and growth.

Key computations: ATV = Gross ÷ Bills · UPT = Qty ÷ Bills · ASP = Gross ÷ Qty · Dis % = Discount ÷ MRP value.

3.2.7 Offer Analysis — Buy 2 Get 30% Off

Source: Buy 2 get 30% off Analysis 9-Sep & 21-Sep.xlsx

Last week vs this week during the offer: qty, gross, discount % by zone; store-wise and product-wise qty growth; bill-level detail.

Key computations: Qty growth = This week – Last week; offer uplift by store and style.

3.2.8 Pujo Offer Analysis & Pujo Working

Source: Pujo Offer Analysis.xlsb · Pujo Working 17th-Aug'26.xlsx

Festive (Pujo) performance for East stores: Pujo'24 vs Pujo'25 gross, qty, bills, ABS, ASP, avg MRP, discount, ATV; bill-value buckets (e.g. 10K & above); plan for Pujo'26 — store × category opening SOH, option count, inward, mix, STR %, projected sale, avg MRP vs ASP.

Key computations: ABS = Qty ÷ Bills · ATV = Gross ÷ Bills · Projection = LY sale × growth assumption; SOH adequacy = plan qty vs opening SOH.

3.2.9 Footfall Aug'25 vs Aug'26

Source: Footfall Aug'25 vs Aug'26.xlsx

Store-wise footfall by grade and zone with growth %, feeding conversion (bills ÷ footfall) once POS bills are joined.

Key computations: Growth % = (Aug'26 – Aug'25) ÷ Aug'25 · Conversion = Bills ÷ Footfall.

3.3 Common dimensions & slicers across the KAZO suite

Zone (Delhi NCR, North, Upper North, East, West, South) · ASM / Cluster · Store & Store Grade (A+, A, B, C, FO) · Location type (EBO / LFR / Online) · L2L status · Category → Sub-category → Option → SKU/Size · Option Grade (A+…E) · Season / Season Group (AW26, SS26, OSM, In-season) · Ageing bucket · Date / Week / Month · Offer / Remarks. Standard measures: Sale Qty, Gross Sale, MRP Value, Discount & Dis %, Bills, ATV, UPT, ABS, ASP, SOH, In-transit, WH Qty, STR %, ROS, Days Cover, Sell-out %, Ach %, Growth %.

3.4 Automation that removes the manual work

- Scheduled pulls of the BC/POS sales, stock and transfer extracts (SFTP / mailbox / API) with the same recognition & validation engine — no more copy-paste of 40k-row dumps into each workbook.
- Single item & store master with versioning; grade changes and MDQ revisions tracked, so L2L and grade-wise views stay consistent across reports.
- Ageing, broken-size, STR and projection routines run nightly; alerts for options crossing ageing thresholds, stores with broken A+/A options, offer weeks under-performing, stores below day target.
- Ask KAZO answers questions across all of the above (e.g. 'which A+ options are broken in Delhi NCR A+ stores?') from the same data.

4. Technical Architecture & Stack

Layer	Technology	Notes
Frontend	React 19, Tailwind CSS, shadcn/ui, Recharts, react-markdown	SPA; role-gated routing; URL-state for shareable report views
Backend / API	Python 3.11, FastAPI (async), Pydantic, Uvicorn	REST under /api; JWT auth; SSE streaming for Ask KAZO
Database	MongoDB (Motor async driver)	Layers: raw_files → raw_rows → std_rows (versioned) → sales → exceptions; masters; audit_logs; indexes on all filter keys
Ingestion engine	pandas, openpyxl/xlrd, rapidfuzz	Deterministic: header detection, alias dictionary, fuzzy matching, learned mappings, 12 validation rules
Analytics	MongoDB aggregation pipelines	Server-side computation of all measures (distinct counts, panels, targets) — no pre-aggregated cubes
AI layer	OpenAI GPT-5.4 via Emergent Universal Key, function-calling	Read-only tools over the reporting DB; grounded system prompt; per-user sessions
Storage	Managed object storage	Original partner files kept immutable; download from batch
Security	bcrypt, JWT (12 h), brute-force lockout, Fernet-encrypted connector secrets, role guards, audit trail	HTTPS; CORS restricted; secrets in environment
Hosting	Kubernetes-managed containers (frontend, API, DB), CI-style deploys	Preview + production environments; horizontal scaling of API
Quality	pytest engine tests, automated UI/API regression suites (7 iterations)	Every release verified against critical flows

Data model (reporting)

secondary_sales: partner, week, month, store_code, store_name, style, size, colour, nsv, qty, mrp, cost, category, sub_category, year, till_date, ltl, margin, area_manager, top_stores, existing_vs_new, city, state, target + pre-computed row components (mrp_val, cost_val, gst_val, gst_reversal, margin_val, margin_nsv, panel_w, ltl_nsv). **sales** (partner weekly files via pipeline) with full batch lineage. Masters: partners, stores/counters, styles/SKUs, targets, stock, margins, aliases, partner_mappings, validation_rules — all loaded from KAZO templates.

5. Computation Standards (extract from Help Center)

Measure	Definition
NSV-yy	Σ NSV where YEAR = yy (2025 / 2026)
CONT %	NSV-26 of row ÷ NSV-26 of all rows in filter
NSV G / QTY G	(value-26 – value-25) ÷ value-25
LTL G	Growth on rows where LTL = 'LTL' (store has sales in the month in both years)
#STR-yy	Distinct stores with NSV > 0 in year
#PANEL-yy	Σ over trading stores of Σ over categories with NSV > 0 of panel weight (configurable per category)
EST / EST G	Σ over (store, week) of max(target), YEAR 2026; EST G = NSV-26 ÷ EST – 1
PROD/WK-yy	NSV-yy ÷ distinct stores ÷ distinct weeks
GST VAL	Σ NSV – NSV ÷ (1 + rate), rate 5 % if unit NSV ≤ 2,500 else 18 %
PRIM BILLING-26	MRP VAL – GST VAL – MARGIN VAL (2025: SECONDARY-PRIMARY column)
REALIZED VALUE	PRIM BILLING – DISC + GST REVERSAL + MARGIN REVERSAL
MARGIN %	(REALIZED VALUE – COST VAL) ÷ REALIZED VALUE
Recognition confidence	0.3 × partner + 0.1 × header + 0.6 × avg column; capped at 45 if a required field is unmapped
Acceptance %	(accepted + warning rows) ÷ total rows
Retail week	Week 1 = 1 Jan → first Sunday, then Monday–Sunday (configurable to KAZO calendar)

6. Hosting, AI Layer & Managed Service (covered by the monthly fee)

- **Hosting**: production Kubernetes environment (frontend, API, MongoDB), TLS, backups (daily, 30-day retention), monitoring and uptime alerts; 99.5 % monthly availability target.
- **AI layer**: GPT-5.4 usage for Ask KAZO up to a fair-use pool of 1,000 questions/month; overage billed at cost.
- **Object storage**: up to 25 GB of original partner files; retention as per KAZO policy (never delete by default).
- **Support**: business-hours L2/L3 support (IST), 4-hour response for P1, monthly health report, security patches and dependency upgrades.
- **Change absorption**: onboarding of new partner formats / alias updates / rule tuning up to 4 hours per month; new masters loads.
- **Continuous improvement**: quarterly review with business owner; minor UI/report tweaks included.

Prepared by **Fundle** for KAZO. We look forward to partnering with KAZO on making retail data trusted, fast and self-service.